

Opinion

Better artificial intelligence does not mean better models of biology

Drew Linsley^{1,*}, Pinyuan Feng², and Thomas Serre¹ 

Deep neural networks (DNNs) once showed increasing alignment with primate perception as they improved on vision benchmarks, raising hopes that advances in artificial intelligence (AI) would naturally yield better models of biological vision. However, we present accumulating evidence that this alignment is now plateauing – and in some cases worsening – as DNNs scale to human or even superhuman accuracy. This divergence between artificial and biological perception may reflect the acquisition of visual strategies distinct from those of primates, and these findings challenge the view that advances in AI will naturally translate to progress in neuroscience. We argue that vision science must chart its own course, developing algorithms grounded in biological visual systems rather than optimizing for internet data.

The transformative promise of deep learning for vision science

One year after AlexNet – the **deep neural network (DNN)** (see [Glossary](#)) that sparked the modern AI revolution – Yamins, DiCarlo, and colleagues [1–3] published a pivotal finding: units in task-optimized DNNs trained for perceptual tasks, such as object recognition, responded to images in similar ways to real neurons in primate **inferotemporal cortex (ITC)**. This discovery transformed vision science and was soon joined by a chorus of similar findings in other domains of neuroscience (e.g., [4,5]). Suddenly, DNNs were not just systems for categorizing dogs and cats; they were now models of biology that could settle long-standing debates about neural computations and their links to perception. For example, the observation that DNNs pretrained for object classification better predicted IT responses provided computational backing to the theory of core object recognition – the idea that primate visual circuitry is optimal for rapid and invariant object recognition [6]. Perhaps most importantly, these newfound capabilities of DNNs implied that optimizing them for computer vision tasks and AI more broadly might uncover the principles that shape biological vision in the process [3,7]. Advances in AI through engineering could now lead the way and biological discovery would follow as a byproduct.

In the decade since AlexNet, DNNs have advanced dramatically, and current state-of-the-art models rival or surpass humans on nearly all vision benchmarks [8,9]. This progress has almost entirely been driven by an exponential scale-up of DNNs, a trend that was catalyzed by the introduction of **Transformers** [10] – attentional architectures that are purpose-built for the parallel processing capabilities of graphics processing units (GPUs) – and the observation that DNN performance on nearly all benchmarks improves predictably as the number of model parameters and amount of data used for training increases [11,12].

Nevertheless, today's state-of-the-art DNNs, including models like those from OpenAI, Anthropic, and Google – estimated to contain trillions of parameters trained on large proportions of the internet – still behave in strange ways; for example, stumbling on simple problems like

Highlights

A surprising feature of early deep neural networks (DNNs) was that they evolved into better models of primate vision as their object-recognition performance improved.

Despite achieving human-level accuracy, today's scaled-up DNNs are worse than ever at predicting primate vision, likely because they rely on visual strategies that are misaligned with primates.

Vision science must chart an independent course, developing deep learning approaches tailored to biological vision rather than following artificial intelligence benchmarks.

¹Brown University, Providence, RI, USA²Columbia University, New York, NY, USA*Correspondence: drew_linsley@brown.edu (D. Linsley).

counting or predicting viewpoints that seem trivial to humans, while excelling at a wide array of more complex ones [13–15]. This chasm between the human-level accuracy and alien-like behavior of today's DNNs raises a critical question for vision science: has the relentless engineering of DNNs to improve on AI benchmarks continued to bring concomitant improvements for modeling brains and behavior, or has it steered these systems away from biological principles altogether?

How deep learning has reshaped vision science

Task-optimized deep learning [16,17] has become extremely popular in computational neuroscience for modeling the brain for multiple reasons. It is the state-of-the-art modeling technique for system identification and predicting primate neural responses to images, an application which is fundamental for neuroprosthetic device development [18–21] and *in silico* studies of function [22–24]. Task-optimization can also be used to glean insights into the principles that sculpt visual circuitry.

For instance, early work found that DNNs trained for object classification predicted neural responses in V4 and ITC better than models trained end-to-end for neural prediction [1,2]. This suggested that primate visual systems are organized around an objective to recognize objects in the world [4,25–39]. Notably, self-supervised learning approaches, which do not explicitly optimize for object classification, can lead to models that are nearly as effective at predicting ITC responses [27,28,40–43]. While task-optimization clearly matters for modeling primate vision, the relationship between specific tasks and neural organization remains an open question.

Task-optimization has also proven to be an effective strategy for modeling human perception. Most prominently, it has been repeatedly demonstrated that the alignment between human and DNN object-recognition decisions in psychophysics experiments increases with model accuracy on computer vision benchmarks [8,9,44–46]. This convergence is more than just a consequence of fewer mistakes; even the patterns of errors produced by these DNNs have come to resemble those of humans [9,47]. Task-optimized DNNs are capable of predicting human judgments across diverse tasks, including tests of local versus global processing biases [48], the perceived semantic similarity of object images [49,50], bottom-up image saliency [51,52], and the 3D properties of objects and scenes [13,53,54].

Others have found that co-training DNNs for object classification and predicting primate neural data can help models better predict human object classification psychophysics [55–57]. However, even with task-optimization, there are still multiple features of human perception that have proven challenging to model without introducing additional mechanisms into the model, such as perceptual phenomena like illusions [58–64] and metamers [65–68] – the case when two materially different stimuli are perceived as identical.

Inspired by the success of engineering benchmarks in advancing AI, the visual neuroscience community has developed its own benchmarks to systematically evaluate the prediction accuracy of models of biological vision [69,70]. A particularly influential initiative is **Brain-Score**, an online benchmark with a living leaderboard that allows anyone to submit and evaluate models on private data from dozens of neural and psychophysics experiments [7,71]. By tracking how DNNs perform on biological benchmarks as they improve on engineering ones, Brain-Score offers essential insights into how advances in AI have translated to biological vision.

The effectiveness of task-optimization has declined for modern DNNs

Despite its early transformative impact, optimizing DNNs for image tasks such as object classification or for learning from image datasets through **self-supervised training** has become an

Glossary

Attribution maps: visualizations that highlight regions of an input image that influence a neural network's predictions, providing insights into model decision-making processes and feature importance for specific classifications.

Brain-Score: an online benchmarking platform for systematically evaluating the accuracy of computational models at predicting neural and behavioral data.

Convolutional neural networks (CNN): DNNs models that apply multiple sets of learnable filters across all positions of an input, with each filter detecting patterns that are needed to solve the objective specified for the model. The same weights are shared across spatial locations. When layers of convolutional filters are stacked, deeper layers apply spatially local filtering to the outputs of previous layers. This process builds hierarchical representations, progressing from simple features to complex patterns.

Deep neural network (DNN): a computational model consisting of multiple processing layers that learn representations of data with multiple levels of abstraction, inspired by the neural networks in the brain. Convolutional neural networks and Transformer networks are special cases.

ImageNet: a large-scale dataset of labeled images spanning thousands of object categories, widely used as a benchmark for training and evaluating object-recognition systems.

Inferotemporal cortex (ITC): a region in the primate visual cortex responsible for high-level visual processing, particularly object recognition and categorization.

Neural alignment: the degree to which neural activity patterns in trained models correlate with or predict neural responses in biological systems when presented with the same stimuli.

Self-supervised training: training models with unlabeled data by creating supervisory signals from the data itself rather than human annotations. In computer vision, this typically involves solving pretext tasks, such as filling in masked image regions. Versions of self-supervised training are used in the training pipelines for large-scale generative models, like chatbots.

Transformer: a deep learning architecture that relies on self-attention mechanisms to process and relate different positions in input sequences,

increasingly unreliable strategy for modeling biological vision. We first observed this decline in the Brain-Score benchmark [32,72]. One of the experimental datasets included in Brain-Score is the same one used by Yamins, DiCarlo, and colleagues in their seminal work [73]. Brain-Score, however, updates those results by including the **neural alignment** of contemporary DNNs with the ITC neurons recorded in that dataset (Figure 1A, results replotted from Brain-Score.com). As newer DNNs have been added over recent years, a striking trend has emerged: DNNs continued to improve as models of ITC until they reached around 70% accuracy (chance is ~0.1%) on **ImageNet**, at which point they explained approximately half of the neural variance. Beyond this point, DNN neural alignment scores have declined even as their ImageNet accuracies have improved. In other words, optimizing DNNs for object classification has proven ineffective for building better models of ITC neurons in this dataset.

This trend of worsening neural alignment in image-classification optimized DNNs appears to be general. Indeed, we found a similar trend in recordings taken from a completely different study [32], in which monkeys were instructed to fixate at different locations across natural images as posterior ITC neurons were being recorded (Figure 1B). In this experiment, we benchmarked the neural alignment of a large, diverse, and highly controlled set of DNNs from the Pytorch Image Model library (TIMM) (<https://github.com/rwightman/pytorch-image-models>), which represented a decade's worth of progress in computer vision and included all of the popular architectures (**convolutional neural networks**, **Vision Transformers**, and hybrids of these two) and pretraining approaches (e.g., first training on ImageNet-21K [74], using self-supervised training approaches like DINO [75], or using multimodal learning approaches with image-language modeling approaches like CLIP [76]) over that period (note that Figure 1B depicts an updated version of the analysis in [32] with models since 2025).

Each of the DNNs was evaluated on two objectives: (i) accuracy on a multilabel ImageNet variant that was benchmarked against humans [77], and (ii) alignment with primate ITC neuron responses to specific regions of natural images [73] using the Brain-Score procedure [7,71]. What made this object classification benchmark unique was that it addressed a key issue in the original ImageNet, where multiple objects in some images had confounded earlier attempts to evaluate standard single-object classification accuracy. It also contained performance from five human participants on 80% of the ImageNet validation set, enabling direct human-to-machine comparisons for object recognition performance.

Many of the DNNs in this set had object classification performance that rivaled or surpassed humans: 30% (111) were human-level (scoring within the human confidence interval) and 1% (4) outperformed humans (the EVA class of models [78]; Figure 1B). At the same time, these human-level DNNs had also devolved into poor models of primate vision. For example, a super-human DNN (EVA-02, released in 2023, only included in our re-analysis) was as aligned with ITC as one from 2019 that was far below human-level (the MobileNetV3 [79] from 2019). We also localized the inflection point at which the previously positive correlation between task performance and neural alignment turned negative at the TResNet – a GPU performance-optimized implementation of the original ResNet [80,81]. Follow-up analyses by us and others, comparing model performance across benchmarks such as the original ImageNet and more challenging variants, have shown similar reversals in neural alignment [72,82].

To understand why modern DNNs have diverged from biological vision, we analyzed how architectural and training factors influence neural alignment across hundreds of models. When we compared older, less-accurate models with newer, more accurate ones (less vs. more accurate than the 'TResNet' DNN inflection point in neural alignment), a striking pattern emerged. For older


initially developed for natural language processing but later adapted for computer vision tasks like image or video processing.

Vision Transformer (ViT): a type of transformer architecture adapted for computer vision tasks that processes images as sequences of patches rather than through traditional convolutional operations.

V4: an intermediate visual area in the ventral stream of the primate visual cortex that is thought to process complex visual features like shape and color, and which serves as a bridge between early visual areas and higher-level object-recognition regions.

(A) Predicting inferotemporal cortex (ITC) image responses



Majaj et al., 2015



Modeling approach

Schrimpf et al., 2020

DNNs uploaded to Brain-Score.com



(B) Predicting image-mapped ITC responses



Arcaro et al., 2020



Modeling approach


Linsley*, Felipe*, et al., 2023

TIMM DNNs



(C) DNN alignment with human feature importance



Linsley et al., 2019



Modeling approach

Fel*, Felipe*, Linsley* et al., 2022

TIMM DNNs



Trends in Cognitive Sciences

Figure 1. Deep neural networks (DNNs) increasingly diverge from biological vision as recognition accuracy improves. (A) DNNs optimized for visual tasks like object recognition are the standard for predicting image-evoked responses of neurons in the primate inferotemporal cortex (ITC) [1]. Units in each layer of an image task-optimized DNN are linearly fit to a training set of recordings and then the best-fitting layer is used to make predictions on independent recordings [7]. While there once was a strong and linear relationship between the object-recognition accuracy of a model and its ability to predict ITC responses to images, the sign of that correlation has flipped as models have surpassed 70% accuracy (chance is ~0.1%) on ImageNet. Each dot represents a different model. Broken lines denote trends of models on the pareto frontier of neural alignment (i.e., best-case biological predictions). The model representing the changepoint for the trends is the VOneGRCNN-52e-full [104], a recurrent convolutional neural network. (B) A similar pattern of results was found in another study of ITC [32,117], which sought to map the firing rates of neurons to different parts of the same image, yielding neural activity 'heatmaps' for images. The initially positive relationship between DNN object classification accuracy and neural predictivity flipped with the release of the 'TresNet' [88,113]. The gray patch denotes human accuracy on this version of ImageNet [74]. (C) One explanation for this flipped trend is that DNNs have learned to base their decisions on visual features that are different from those of humans (<https://github.com/rwightman/pytorch-image-models>) [88,114,115]. In this experiment [88], using data from [87], visual feature importance maps were derived from human participants who identified parts of images they considered diagnostic for recognition. DNN feature importance maps were then derived using standard

(Figure legend continued at the bottom of the next page.)

generations of models, better object-recognition and larger training datasets reliably improved neural alignment, confirming what the field had long believed [1,2,7,71] (Figure 2A). However, for today's high-performing models, this relationship has broken down entirely. Neither architectural choices (convolutional networks vs. transformers), model scale (parameters, layers, input size), nor training data (ImageNet alone vs. massive internet-scale datasets like ImageNet-21K or LAION-5B [83]) could explain which models better matched the properties of primate neurons. This disconnect is particularly striking given that these same factors – especially scale – are precisely what enabled DNNs to reach human-level performance. In other words, the computational strategies driving AI forward appear fundamentally different from those underlying biological vision.

Why does task-optimization now lead to worse models of biological vision?

One possible explanation for the declining effectiveness of task-optimized deep learning is that DNNs have begun to discover visual strategies that are challenging for biological visual systems to exploit as they have been scaled up. Early hints of this problem came from Ullman and colleagues [84], who found that early DNNs, such as AlexNet and VGG [85], reached their object-recognition decisions differently than humans. Humans exhibited an all-or-nothing reliance on diagnostic features, and the excision of a small number of pixels in an image was sufficient to render it no longer recognizable to a human observer. DNNs did not show the same tendency. These results revealed a new and surprising dimension for comparing human and DNN vision. At the same time, the study used a very small dataset and was limited in multiple ways [86], leaving open the question of whether these differences between humans and DNNs were mere experimental quirks or more fundamental.


To answer this question, we developed ClickMe, a scalable, gamified platform for collecting high-resolution feature importance maps for large natural image datasets [87–89]. In ClickMe, human participants highlight image regions that they find diagnostic for recognition, and while they do this, a DNN tries to recognize the object based only on dilated versions of those regions (Figure 1C). The faster participants can help the DNN recognize the object, the more likely they are to win a prize. The resulting maps were highly consistent across individuals and highlighted features significantly more diagnostic for human object classification than bottom-up saliency [87,88]. The diagnosticity of features in ClickMe maps was further validated in independent psychophysics experiments, showing that another group of participants was significantly better at classifying objects when using features highlighted in ClickMe than in bottom-up saliency maps. We collected these maps for a large proportion of ImageNet and compared their alignment with analogous feature-**attribution maps** from DNNs [90].

This experiment revealed a growing gap between artificial and human vision: as DNNs have advanced beyond human accuracy, the features they have learned for recognition have become progressively more distinct from those used by humans. For example, an 'adversarially-trained Wide ResNet50' achieved the highest alignment with humans, capturing ~50% of the variance in human feature importance maps. More accurate models showed precipitously declining human alignment, increasingly relying on features that differed from those used by humans – such as background textures, global statistics, and even watermark labels (Figure 3) [32,88,91]. Statistical analyses of DNN design factors confirmed this pattern (Figure 2B). While training data, architecture, and scale significantly predicted human feature alignment in older, less-accurate models (before the adversarially-trained Wide ResNet in Figure 1C), these


artificial intelligence interpretability methods based on the model's decision gradient. The alignment between humans and DNNs was measured by computing the correlation between these maps. The changepoint model is an 'Adversarially-trained Wide ResNet50' [116]. See [1,7,32,73,87,88,104,113–117]. Abbreviations: CNN, convolutional neural network; ViT, Vision Transformer.

(A) Predicting image-mapped ITC responses

Experimental data from Arcaro et al., 2020

**(B) DNN alignment with human feature importance**

Experimental data from Linsley et al., 2020



Trends in Cognitive Sciences

Figure 2. Deep neural network (DNN) design factors and biological alignment. We used analysis of variance (ANOVA) models to estimate how DNN design factors predict biological alignment on two experimental datasets: (A) predicting image-mapped inferotemporal cortex (ITC) responses [32,117], and (B) DNN alignment with human feature importance [87]. Our analysis estimated the design factor and alignment relationship for two groups of DNNs per experiment: those that fell before versus after the DNN ‘change points’ identified in Figure 1 and the main text, in which the positive relationship between ImageNet accuracy and biological alignment flipped to a negative one. The dependent variable in these ANOVAs was biological alignment and the independent variables described design factors for the models that produced these scores: accuracy on multilabel ImageNet (continuous), if the model was pretrained on an internet-scale image or multimodal dataset (binary factor: ImageNet-only versus ImageNet-21K/LAION-5B pretraining), if it was a convolutional neural network (CNN) or Vision Transformer (ViT) (binary), the number of layers (continuous), image resolution (continuous), and model parameter count (continuous). All independent variables were normalized to zero mean and unit variance across each experiment to make their effect sizes comparable. We fit separate ANOVAs for DNNs before versus after the change points in each experiment. We extracted parameter estimates for each independent variable and their robust standard errors from each ANOVA, which we refer to as ‘Effect size’ in the figure. In both experiments, the alignment of older and less accurate DNNs (before the change point) improved with internet-scale pretraining and better ImageNet accuracy. However, for newer, more accurate DNNs (after the change point), these relationships broke down entirely – neither architectural choices, model scale, nor training data could explain which models better matched biological vision. We used Wald tests to evaluate if the difference between parameter estimates for DNNs before versus after the DNN change points were significantly different: $t = (\beta_{\text{pre}} - \beta_{\text{post}}) / \sqrt{(\text{SE}_{\text{pre}}^2 + \text{SE}_{\text{post}}^2)}$. Significance markers (*, $P < 0.05$; ***, $P < 0.001$) denote differences between the two groups, highlighting how the computational strategies driving artificial intelligence progress have diverged from those underlying biological vision. See [32,87,117].

relationships disappeared entirely in more accurate models – mirroring the breakdown observed for neural predictivity (Figure 2A). These results underscore the fundamental shift in DNN optimization strategies that has taken place over recent years: the scale-up of DNNs to human-level performance has come at the cost of modeling biology. Today’s state-of-the-art DNNs achieve their performance through a mixture of human-like and distinctly non-biological strategies. While these models retain some alignment with biological vision, they increasingly supplement evolutionarily-honed mechanisms with alternative – and often inscrutable – solutions that may be highly valuable for non-ethological tasks such as biomedical image analysis [92,93], but are less informative for understanding brains and behavior.

If task-optimized DNNs no longer produce better biological models, what will?


The magnificent promise of task-optimized deep learning was in its suggestion that all we needed was more engineering, including data and compute to precisely model biological vision. The now-

**Human
ClickMe**

Linsley et al., 2020

**Standard
ViT small**

Dosovitskiy et al., 2021

**Harmonized
ViT small**

Fel* et al., 2022


**Trends In Cognitive Sciences**

Figure 3. Deep neural networks (DNNs) can be trained to adopt human-like visual strategies. (Left) ClickMe [87] maps identify pixels humans believe are important for recognizing these object images. (Middle and right) Decision attribution maps for the same images from a Vision Transformer (ViT) [118] and harmonized ViT [88] highlight pixels each model uses for recognition. Without harmonization, ViTs often rely on contextual features like watermarks and backgrounds that are less important to humans. See [87,88,118].

waning effectiveness of this approach presents an opportunity for vision scientists to rethink how we model brains. The first question we must ask is whether DNNs are the right modeling framework to begin with [65].

We sought to answer this question by using a supervised learning approach, which we refer to as ‘harmonization’, that directly optimizes for biological alignment [88]. If DNNs are a reasonable modeling framework going forward, then using the harmonization method to guide them toward object decision attribution maps (Figure 1C) that are as close as possible to human feature importance maps (e.g., collected with ClickMe) during ImageNet training would validate that possibility. We found that harmonized DNNs not only relied on human-like visual features for object

classification (Figure 3), but they also learned human-like sensitivity to adversarial attacks [91]. Furthermore, their image representations were significantly more aligned with ITC neurons than those of standard ‘unharmonized’ models [32].

While there have been recent calls for shifting the brain modeling paradigm away from deep learning [65], our harmonization experiments indicate that DNNs remain a viable approach for building predictive models of biological vision [94]. The problems we observe with DNNs trained on internet data, which are increasingly misaligned with biological vision (Figure 1), reflect a mismatch between the training objectives and the data. Co-training with biological data through harmonization can address this.

However, the breakthrough of task-optimized deep learning was that it could reveal principles and test mechanistic-level hypotheses through the principle that by optimizing for object recognition alone, we could naturally discover the computational strategies evolution selected for. Approaches like harmonization, by contrast, enforce biological alignment without providing insight into the origins of vision. Identifying these atomic components will require renewed focus on searching for data diets, behavioral objectives, and routines for training DNNs that guide them to perceive and encode the world like humans from the outset.

Rethinking task-optimization for biological discovery in the era of internet-scale compute

Computational models in vision science have long operated at smaller scales than frontier AI systems. This reflects at least three factors: the academic roots of vision science, a well-founded desire to rely on reductionism for testing mechanistic hypotheses, and the fact that most existing efforts to incorporate biological mechanisms into DNNs do not easily scale. For example, many biologically-inspired DNNs feature architectural constraints, such as recurrence [33,61,93,95–97] and different forms of feedback [59,60,98], which are known to play key roles in primate vision [99–102]. While these approaches are valuable for testing the impact of different mechanisms on complex behaviors [59], recurrent circuits are typically awkward fits for GPUs because the algorithms used to train them [96,100] make it difficult to scale to large datasets.

These scaling challenges are particularly problematic because biological alignment may emerge only from the right combination of data, objective, and circuit constraints when models are trained at very large scales. Moreover, it is possible that some of the computational strategies of biologically-inspired DNNs could be rediscovered by an easy-to-scale DNN trained with the ‘right’ data and objective [101,102], raising questions about the value of inducing these constraints in the first place. Thus, it may be that the most effective approach to reverse-engineering vision with DNNs is to bypass architecture questions altogether: begin with the same unbiased and scalable DNN architectures that are popular in computer vision (e.g., Vision Transformers), and focus efforts on searching for the data diets and training routines that cause this system to exhibit similar representations and behavior as biological vision. Once we identify data and training routines that successfully align task-optimized models with biology, we can then build a comprehensive explanation of biology [3,103] by incorporating other biological features that may require innate constraints or physical embodiment to emerge, such as topography [19–21] or neuronal noise [104].

We propose that better models of primate vision will emerge from training DNNs with biologically-grounded data diets and learning principles (see Outstanding questions). While others have made similar calls before [3,105–107], the field has yet to seriously pursue this direction at scales that may be infeasible for academic research, yet also necessary to observe biological features

emerge from models. The exponential growth in computational resources over recent years now makes it feasible to move beyond the convenient but biologically implausible paradigm of training on large databases of static images [108,109]. Instead, we can explore how biological alignment emerges as models learn from temporally structured and embodied experiences – the same kinds of sensory streams and environmental pressures that actually shape developing visual systems.

Modern simulation environments and video generation techniques [110–112] have reached sufficient fidelity to create naturalistic training data on demand at scale, while advances in self-supervised training have sketched out spaces of behavioral objectives that can be systematically evaluated with these data. The key insight is that we should not be searching for a single ‘correct’ objective, but rather mapping the manifold of developmental and ecological constraints that naturally produce human-like vision. Only through this systematic exploration can we build models that predict biological vision and reveal its underlying principles, ultimately transforming our understanding of why we see the world the way we do in the process.

Concluding remarks

DNNs once demonstrated increasing alignment with primate vision as they improved on object-recognition benchmarks, suggesting that advances in artificial vision would naturally yield better models of biological vision. However, this alignment has plateaued – and in some cases worsened – as DNNs have scaled to human-level accuracy, likely because modern systems rely on visual strategies that diverge from those of biological systems. These findings underscore the need for vision science to chart its own course in computational modeling that is grounded in biological principles rather than purely optimized for internet data.

Acknowledgments

We are grateful to Dan Yamins, Martin Schrimpf, and Lakshmi Govindarajan for their help with editing and refining this paper. This work is supported by NSF (IIS-2402875), ONR (N00014-24-1-2026), and the ANR-3IA Artificial and Natural Intelligence Toulouse Institute (ANR-19-PI3A-0004).

Declaration of interests

The authors declare no competing interests.

References

1. Yamins, D.L.K. *et al.* (2014) Performance-optimized hierarchical models predict neural responses. higher visual cortex. *Proc. Natl. Acad. Sci. U. S. A.* 111, 8619–8624
2. Yamins, D.L. *et al.* (2013) Hierarchical modular optimization of convolutional networks achieves representations similar to Macaque IT and human ventral stream. In *Proceedings of the 27th International Conference on Neural Information Processing Systems*, pp. 3093–3101, NeurIPS
3. Yamins, D.L.K. and DiCarlo, J.J. (2016) Using goal-driven deep learning models to understand sensory cortex. *Nat. Neurosci.* 19, 356–365
4. Khaligh-Razavi, S.-M. and Kriegeskorte, N. (2014) Deep supervised, but not unsupervised, models may explain IT cortical representation. *PLoS Comput. Biol.* 10, e1003915
5. Kriegeskorte, N. *et al.* (2015) Deep neural networks: a new framework for modeling biological vision and brain information processing. *Annu. Rev. Vis. Sci.* 1, 417–446
6. DiCarlo, J.J. *et al.* (2012) How does the brain solve visual object recognition? *Neuron* 73, 415–434
7. Schrimpf, M. *et al.* (2020) Integrative benchmarking to advance neurally mechanistic models of human intelligence. *Neuron* 108, 413–423
8. Dehghani, M. *et al.* (2023) Scaling vision transformers to 22 billion parameters. In *Proceedings of Machine Learning Research* (2023), pp. 7480–7512
9. Geirhos, R. *et al.* (2021) Partial success in closing the gap between human and machine vision. In *Advances in Neural Information Processing Systems* (34) (Ranzato, M. *et al.*, eds), NeurIPS
10. Vaswani, A. *et al.* (2017) Attention is all you need. In *Proceedings of the 31st Conference on Neural Information Processing Systems*, NeurIPS
11. Kaplan, J. *et al.* (2020) Scaling laws for neural language models. *arXiv* Published online January 23, 2020. <https://doi.org/10.48550/arXiv.2001.08361>
12. Zhai, X. *et al.* (2022) Scaling vision transformers. In *Proceedings of IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE
13. Linsley, D. *et al.* (2024) The 3D-PC: a benchmark for visual perspective taking in humans and machines. In *The Thirteenth International Conference on Learning Representations*, ICLR
14. McCoy, R.T. *et al.* (2024) Embers of autoregression show how large language models are shaped by the problem they are trained to solve. *Proc. Natl. Acad. Sci. U. S. A.* 121, e2322420121
15. Wu, Z. *et al.* (2023) Reasoning or reciting? Exploring the capabilities and limitations of language models through counterfactual tasks. *arXiv* Published online March 28, 2024. <https://doi.org/10.48550/arXiv.2307.02477>

Outstanding questions

How can we create more biologically realistic training regimens for DNNs? Today's DNNs are trained on massive image collections that bear little resemblance to the temporally structured and rich, multimodal experiences that actually shape biological visual systems. How can we bridge this gap? Training DNNs within ecologically valid and dynamic environments using state-of-the-art computer graphics methods.

What learning principles will align DNNs with biological vision? The objective functions used in today's DNNs are designed with static image data in mind. What happens when we move our models to dynamic and embodied data collection? It is essential for the field to systematically explore objectives and data curricula that might cause DNNs to learn more human-like visual representations.

What role will biological constraints play in the search for data diets and learning principles? It is challenging to imagine how certain features of the brain – such as topographical organization and neural noise – could emerge without being explicitly induced. Determining which constraints are fundamental principles and which are byproducts of biological implementation will be crucial for developing models that accurately predict both neural responses and behavior.

16. Richards, B.A. *et al.* (2019) A deep learning framework for neuroscience. *Nat. Neurosci.* 22, 1761–1770
17. Kell, A.J.E. *et al.* (2018) A task-optimized neural network replicates human auditory behavior, predicts brain responses, and reveals a cortical processing hierarchy. *Neuron* 98, 630–644
18. Schiatti, L. *et al.* (2023) Modeling visual impairments with artificial neural networks: a review. In *Proceedings of the 2023 IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, pp. 1987–1999, IEEE
19. Rath, N. *et al.* (2025) TopoLM: brain-like spatio-functional organization in a topographic language model. In *The Thirteenth International Conference on Learning Representations, ICLR*
20. Margalit, E. *et al.* (2024) A unifying framework for functional organization in early and higher ventral visual cortex. *Neuron* 112, 2435–2451
21. Schrimpf, M. *et al.* (2024) Do topographic deep ANN models of the primate ventral stream predict the perceptual effects of direct IT cortical interventions? *bioRxiv* Published online January 9, 2024. <https://doi.org/10.1101/2024.01.09.572970>
22. Guo, C. *et al.* (2022) Adversarially trained neural representations are already as robust as biological neural representations. *Proc. Mach. Learn. Res.* 162, 8072–8081
23. Bashivan, P. *et al.* (2019) Neural population control via deep image synthesis. *Science* 364, 6439
24. Kar, K. *et al.* (2019) Evidence that recurrent circuits are critical to the ventral stream's execution of core object recognition behavior. *Nat. Neurosci.* 22, 974–983
25. Cadena, S.A. *et al.* (2024) Diverse task-driven modeling of macaque V4 reveals functional specialization towards semantic tasks. *PLoS Comput. Biol.* 20, e1012056
26. Willeke, K.F. *et al.* (2022) The Sensorium competition on predicting large-scale mouse primary visual cortex activity. *arXiv* Published online June 17, 2022. <https://doi.org/10.48550/arXiv.2206.08666>
27. Konkle, T. and Alvarez, G.A. (2022) A self-supervised domain-general learning framework for human ventral stream representation. *Nat. Commun.* 13, 491
28. Conwell, C. *et al.* (2024) A large-scale examination of inductive biases shaping high-level visual representation in brains and machines. *Nat. Commun.* 15, 9383
29. Willeke, K.F. *et al.* (2023) Deep learning-driven characterization of single cell tuning in primate visual area V4 unveils topological organization. *bioRxiv* Published online June 17, 2022. <https://doi.org/10.1101/2023.05.12.540591>
30. Cadena, S.A. *et al.* (2019) How well do deep neural networks trained on object recognition characterize the mouse visual system? In *Real Neurons & Hidden Units: Future Directions at the Intersection of Neuroscience and Artificial Intelligence*, NeurIPS
31. Safarani, S. *et al.* (2021) Towards robust vision by multi-task learning on monkey visual cortex. In *Proceedings of the 35th International Conference on Neural Information Processing Systems*, pp. 739–751, NeurIPS
32. Linsley, D. *et al.* (2023) Performance-optimized deep neural networks are evolving into worse models of inferotemporal visual cortex. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, pp. 28873–28891, NeurIPS
33. Kietzmann, T.C. *et al.* (2019) Recurrence is required to capture the representational dynamics of the human visual system. *Proc. Natl. Acad. Sci. U. S. A.* 116, 21854–21863
34. Mehrer, J. *et al.* (2021) An ecologically motivated image dataset for deep learning yields better models of human vision. *Proc. Natl. Acad. Sci. U. S. A.* 118, e2011417118
35. Cadieu, C.F. *et al.* (2014) Deep neural networks rival the representation of primate IT cortex for core visual object recognition. *PLoS Comput. Biol.* 10, e1003963
36. Hong, H. *et al.* (2016) Explicit information for category-orthogonal object properties increases along the ventral stream. *Nat. Neurosci.* 19, 613–622
37. Cichy, R.M. *et al.* (2016) Comparison of deep neural networks to spatio-temporal cortical dynamics of human visual object recognition reveals hierarchical correspondence. *Sci. Rep.* 6, 27755
38. Dobs, K. *et al.* (2023) Behavioral signatures of face perception emerge in deep neural networks optimized for face recognition. *Proc. Natl. Acad. Sci. U. S. A.* 120, e2220642120
39. Dobs, K. *et al.* (2022) Brain-like functional specialization emerges spontaneously in deep neural networks. *Sci. Adv.* 8, eabl8913
40. Zhuang, C. *et al.* (2021) Unsupervised neural network models of the ventral visual stream. *Proc. Natl. Acad. Sci. U. S. A.* 118, e2014196118
41. Mineault, P. *et al.* (2021) Your head is there to move you around: goal-driven models of the primate dorsal pathway. In *Proceedings of the 35th International Conference on Neural Information Processing Systems*, pp. 28757–28771, NeurIPS
42. Cueva, C.J. and Wei, X.X. (2018) Emergence of grid-like representations by training recurrent neural networks to perform spatial localization. *arXiv* Published online March 21, 2018. <https://doi.org/10.48550/arXiv.1803.07770>
43. Storrs, K.R. *et al.* (2021) Diverse deep neural networks all predict human inferior temporal cortex well, after training and fitting. *J. Cogn. Neurosci.* 33, 2044–2064
44. Rajalingham, R. *et al.* (2015) Comparison of object recognition behavior in human and monkey. *J. Neurosci.* 35, 12127–12136
45. Rajalingham, R. *et al.* (2018) Large-scale, high-resolution comparison of the core visual object recognition behavior of humans, monkeys, and state-of-the-art deep artificial neural networks. *J. Neurosci.* 38, 7255–7269
46. Geirhos, R. *et al.* (2018) Generalisation in humans and deep neural networks. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, pp. 7549–7561, NeurIPS
47. Geirhos, R. *et al.* (2020) Beyond accuracy: quantifying trial-by-trial behaviour of CNNs and humans by measuring error consistency. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, pp. 13890–13902, NeurIPS
48. Hermann, K. *et al.* (2020) The origins and prevalence of texture bias in convolutional neural networks. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, pp. 19000–19015, NeurIPS
49. Battleday, R.M. *et al.* (2021) From convolutional neural networks to models of higher-level cognition (and back again). *Ann. N. Y. Acad. Sci.* 1505, 55–78
50. Mahner, F.P. *et al.* (2025) Dimensions underlying the representational alignment of deep neural networks with humans. *Nat. Mach. Intell.* 7, 848–859
51. Kümmerer, M. *et al.* (2022) DeepGaze III: modeling free-viewing human scanpaths with deep learning. *J. Vis.* 22, 7
52. Kümmerer, M. *et al.* (2017) DeepGaze II: reading fixations from deep features trained on object recognition. *J. Vis.* 17, 1147
53. Danier, D. *et al.* (2025) DepthCues: evaluating monocular depth perception in large vision models. In *Proceedings IEEE/CVF Conference on Computer Vision and Pattern Recognition*, IEEE
54. Bonnen, T. *et al.* (2024) Evaluating multiview object consistency in humans and image models. In *Proceedings of the 38th International Conference on Neural Information Processing Systems*, pp. 43533–43548, NeurIPS
55. Dapello, J. *et al.* (2022) Aligning model and macaque inferior temporal cortex representations improves model-to-human behavioral alignment and adversarial robustness. In *The Eleventh International Conference on Learning Representations, ICLR*
56. Seeliger, K. *et al.* (2021) End-to-end neural system identification with neural information flow. *PLoS Comput. Biol.* 17, e1008558
57. St-Yves, G. *et al.* (2023) Brain-optimized deep neural network models of human visual areas learn non-hierarchical representations. *Nat. Commun.* 14, 3329
58. Malhotra, G. *et al.* (2022) Feature blindness: a challenge for understanding and modelling visual object recognition. *PLoS Comput. Biol.* 18, e1009572
59. Linsley, D. *et al.* (2019) Recurrent neural circuits for contour detection. In *International Conference on Learning Representations, ICLR*
60. Kim, J. *et al.* (2020) Disentangling neural mechanisms for perceptual grouping. In *International Conference on Learning Representations, ICLR*

61. Linsley, D. *et al.* (2018) Learning long-range spatial dependencies with horizontal gated recurrent units. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, pp. 152–164, NeurIPS
62. Linsley, D. *et al.* (2021) Tracking without re-recognition in humans and machines. In *Proceedings of the 35th International Conference on Neural Information Processing Systems*, pp. 19473–19486
63. Muzellec, S. *et al.* (2024) Tracking objects that change in appearance with phase synchrony. In *International Conference on Learning Representations*, ICLR
64. Doerig, A. *et al.* (2020) Crowding reveals fundamental differences in local vs. global processing in humans and machines. *Vis. Res.* 167, 39–45
65. Bowers, J.S. *et al.* (2022) Deep problems with neural network models of human vision. *Behav. Brain Sci.* 46, e385
66. Feather, J. *et al.* (2025) Brain-model evaluations need the NeuroAI turing test. *arXiv* Published online February 22, 2025. <https://doi.org/10.48550/arXiv.2502.16238>
67. Feather, J. *et al.* (2019) Metamers of neural networks reveal divergence from human perceptual systems. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, pp. 10078–10089, NeurIPS
68. Feather, J. *et al.* (2023) Model metamers reveal divergent invariances between biological and artificial neural networks. *Nat. Neurosci.* 26, 2017–2034
69. Cichy, R.M. *et al.* (2019) The algonauts project. *Nat. Mach. Intell.* 1, 613
70. Hebart, M.N. *et al.* (2023) THINGS-data, a multimodal collection of large-scale datasets for investigating object representations in human brain and behavior. *eLife* 12, e82580
71. Schrimpf, M. *et al.* (2018) Brain-Score: which artificial neural network for object recognition is most brain-like? *bioRxiv* Published online January 2, 2020. <https://doi.org/10.1101/407007>
72. Gokce, A. and Schrimpf, M. (2024) Scaling laws for task-optimized models of the primate visual ventral stream. In *Forty-second International Conference on Machine Learning*, ICML
73. Majaj, N.J. *et al.* (2015) Simple learned weighted sums of inferior temporal neuronal firing rates accurately predict human core object recognition performance. *J. Neurosci.* 35, 13402–13418
74. Ridnik, T. *et al.* (2021) ImageNet-21K pretraining for the masses. *arXiv* Published online August 5, 2021. <https://doi.org/10.48550/arXiv.2104.10972>
75. Oquab, M. *et al.* (2024) DinoV2: learning robust visual features without supervision. *Trans. Mach. Learn. Res.* Published online January 11, 2024. <https://openreview.net/forum?id=a68SU6zFt>
76. Radford, A. *et al.* (2021) Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, ICML
77. Shankar, V. *et al.* (2020) Evaluating machine accuracy on ImageNet. In *Proceedings of the 37th International Conference on Machine Learning*, ICML
78. Fang, Y. *et al.* (2023) EVA-02: a visual representation for neon genesis. *Image Vis. Comput.* 149, 105171
79. Howard, A. *et al.* (2019) Searching for MobileNetV3. In *Proceedings IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE
80. He, K. *et al.* (2016) Deep residual learning for image recognition. In *IEEE Conference on Computer Vision and Pattern Recognition*, IEEE
81. He, K. *et al.* (2016) Identity mappings in deep residual networks. In *European Conference on Computer Vision* (Leibe, B. *et al.*, eds), pp. 630–645, Springer
82. Nonaka, S. *et al.* (2021) Brain hierarchy score: which deep neural networks are hierarchically brain-like? *iScience* 24, 103013
83. Schuhmann, C. *et al.* (2022) LAION-5B: an open large-scale dataset for training next generation image-text models. In *Advances in Neural Information Processing Systems 35, Datasets and Benchmarks Track* (Koyejo, S. *et al.*, eds), NeurIPS
84. Ullman, S. *et al.* (2016) Atoms of recognition in human and computer vision. *Proc. Natl. Acad. Sci. U. S. A.* 113, 2744–2749
85. Simonyan, K. and Zisserman, A. (2015) Very deep convolutional networks for large-scale image recognition. In *International Conference on Learning Representations*, ICLR
86. Funke, C.M. *et al.* (2021) Five points to check when comparing visual perception in humans and machines. *J. Vis.* 21, 16
87. Linsley, D. *et al.* (2019) Learning what and where to attend. In *Proceedings of the International Conference on Learning Representation*, ICLR
88. Fel, T. *et al.* (2022) Harmonizing the object recognition strategies of deep neural networks with humans. *Adv. Neural Inf. Proces. Syst.* 35, 9432–9446
89. Linsley, D. *et al.* (2017) What are the visual features underlying human versus machine vision? In *2017 IEEE International Conference on Computer Vision Workshops (ICCVW)*, IEEE
90. Simonyan, K. *et al.* (2013) Deep inside convolutional networks: visualising image classification models and saliency maps. In *International Conference on Learning Representations*, ICLR
91. Linsley, D. *et al.* (2023) Adversarial alignment: breaking the trade-off between the strength of an attack and its relevance to human perception. *arXiv* Published online June 5, 2023. <https://doi.org/10.48550/arXiv.2306.03229>
92. Linsley, J.W. *et al.* (2021) Superhuman cell death detection with biomarker-optimized neural networks. *Sci. Adv.* 7, eabf8142
93. Nayebi, A. *et al.* (2018) Task-driven convolutional recurrent models of the visual system. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, NeurIPS
94. Linsley, D. and Serre, T. (2023) Fixing the problems of deep neural networks will require better training data and learning algorithms. *Behav. Brain Sci.* 46, e400
95. Azeglio, S. *et al.* (2022) Improving neural predictivity in the visual cortex with gated recurrent connections. *arXiv* Published online March 22, 2022. <https://doi.org/10.48550/arXiv.2203.11910>
96. Linsley, D. *et al.* (2020) Stable and expressive recurrent vision models. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, pp. 10456–10467, NeurIPS
97. Spoerer, C.J. *et al.* (2020) Recurrent neural networks can explain flexible trading of speed and accuracy in biological vision. *PLoS Comput. Biol.* 16, e1008215
98. Roelfsema, P.R. *et al.* (2000) The implementation of visual routines. *Vis. Res.* 40, 1385–1411
99. Gilbert, C.D. and Li, W. (2013) Top-down influences on visual processing. *Nat. Rev. Neurosci.* 14, 350–363
100. van Bergen, R.S. and Kriegeskorte, N. (2020) Going in circles is the way forward: the role of recurrence in visual inference. *Curr. Opin. Neurobiol.* 65, 176–193
101. Huh, M. *et al.* (2024) The platonic representation hypothesis. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 20617–20642, ICML
102. Liao, Q. and Poggio, T. (2016) Bridging the gaps between residual learning, recurrent neural networks and visual cortex. *arXiv* Published online April 13, 2016. <https://doi.org/10.48550/arXiv.1604.03640>
103. Doerig, A. *et al.* (2023) The neuroconnectionist research programme. *Nat. Rev. Neurosci.* 24, 431–450
104. Dapello, J. *et al.* (2020) Simulating a primary visual cortex at the front of CNNs improves robustness to image perturbations. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, pp. 13073–13087, NeurIPS
105. Kietzmann, T.C. *et al.* (2019) Deep neural networks in computational neuroscience. In *Oxford Research Encyclopedia of Neuroscience*, Oxford University Press
106. Orhan, A.E. and Lake, B.M. (2024) Learning high-level visual representations from a child's perspective without strong inductive biases. *Nat. Mach. Intell.* 6, 271–283
107. Vogelsang, L. *et al.* (2018) Potential downside of high initial visual acuity. *Proc. Natl. Acad. Sci. U. S. A.* 115, 11333–11338
108. Zador, A. *et al.* (2023) Catalyzing next-generation artificial intelligence through NeuroAI. *Nat. Commun.* 14, 1597
109. Warren, W.H. (2021) Information is where you find it: perception as an ecologically well-posed problem. *i-Perception* 12, 2
110. Mildenhall, B. *et al.* (2020) NeRF: representing scenes as neural radiance fields for view synthesis. In *European Conference on Computer Vision* (Leibe, B. *et al.*, eds), pp. 405–421, Springer

111. Kerbl, B. *et al.* (2023) 3D Gaussian splatting for real-time radiance field rendering. *ACM Trans. Graph.* 42, 1–14
112. Wu, G. *et al.* (2023) 4D Gaussian splatting for real-time dynamic scene rendering. In *IEEE Conference on Computer Vision and Pattern Recognition*, IEEE
113. Ridnik, T. *et al.* (2021) TResNet: high performance GPU-dedicated architecture. In *Proceedings of IEEE Winter Conference on Applied Computer Vision (WACV)*, IEEE
114. Geirhos, R. *et al.* (2020) Shortcut learning in deep neural networks. *Nat. Mach. Intell.* 2, 665–673
115. Shahamatdar, S. *et al.* (2024) Deceptive learning in histopathology. *Histopathology* 85, 116–132
116. Xie, C. *et al.* (2019) Feature denoising for improving adversarial robustness. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE
117. Arcaro, M.J. *et al.* (2020) The neurons that mistook a hat for a face. *eLife* 9, e53798
118. Dosovitskiy, A. *et al.* (2021) An image is worth 16x16 words: transformers for image recognition at scale. In *International Conference on Learning Representation*, ICLR